

Extended TD3_vanilla training: training-episode reward (noisy episodes over random scenarios -- NOT a convergence signal on its own; see f13 for the deterministic validation curve)

Training-episode reward
(SqTrError_TrPenalty_UserIncentives, with exploration noise)

