All models on the final evaluation grid (station_v0_bogota, 50 scenario seeds x weekday/weekend = 100 cells per model)
Point = mean; bar = 95% CI over the 50 seed-level means (ENS_rel: cluster-bootstrap CI, S5.7)

Heuristics and control
AFAP
Round Robin
Random (control)

Information-advantaged (non-causal)
MPC (tracking)
MPC (energy-max)
Oracle (tracking-only)
Oracle (balanced)

TD3 original, 60k (trained pre-setpoint-fix)
TD3 (seed 100)
TD3 (seed 101)
TD3 (seed 102)
TD3-TrackingOnly (seed 100)
TD3-TrackingOnly (seed 101)
TD3-TrackingOnly (seed 102)

TD3 vanilla, new run @60k
TD3 new run @60k (seed 100)
TD3 new run @60k (seed 101)
TD3 new run @60k (seed 102)

TD3 vanilla, extended (primary)
TD3 extended, primary (seed 100)
TD3 extended, primary (seed 101)
TD3 extended, primary (seed 102)

TD3 vanilla, extended (last)
TD3 extended, last (seed 100)
TD3 extended, last (seed 101)
TD3 extended, last (seed 102)

15% target

10,000 20,000 30,000 40,000 50,000 60,000
Tracking error
(sum of squared kW deviations; lower is better)

0 5 10 15 20
Transformer overload
(kW x 15-min steps per day; / 4 = kWh)

96 97 98 99 100
Average user satisfaction
(%; higher is better; target > 90%)

0 5 10 15 20
Energy not served vs. AFAP, ENS_rel
(%; target: 95% CI upper bound < 15%)

Round Robin (recommended strategy, S5.8) -- reference line
Declared target